

# Memory Footprint of a Java Process

Zhihu: tkanng

GitHub: tkanng

Email: tkanng@gmail.com

# Typical Questions

- Why is Java using much more memory than heap size?
- How to keep the memory used by a Java process under control?
- How to detect the memory usage of off-heap?
- Why does NMT (Native Memory Tracking) report more / less committed memory than the linux process resident set size?
- .....

# Memory Footprint of a Java Process

- Linux Memory Management
- Memory Footprint of a Java process
- Tools to detect memory usage of a Java process
- Examples

# **Linux Memory Management**

# Linux Memory Management

- RSS(RES): Resident Set Size
- VIRT: Virtual memory used by the task. (It matters for 32-bit OS, cause it's maximum virtual memory is just 4G)

PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME
10	-10	128M	17252	14804	R	1.3	0.8	2h18:
10	-10	128M	17252	14804	S	0.7	0.8	35:03.
20	0	26092	3556	2784	R	0.7	0.2	0:00.
10	-10	128M	17252	14804	R	0.7	0.8	12:22.
10	-10	128M	17252	14804	R	0.7	0.8	12:53.
20	0	749M	40228	0	S	0.7	2.0	11:06.
10	-10	128M	17252	14804	S	0.0	0.8	4:37.
10	-10	128M	17252	14804	S	0.0	0.8	8:27.
20	0	1082M	116M	0	S	0.0	5.8	20:30.
20	0	40868	4648	4048	S	0.0	0.2	2:24.
10	-10	32392	6276	5524	S	0.0	0.3	1:49.
10	-10	128M	17252	14804	S	0.0	0.8	7:09.
10	-10	427M	5696	5184	S	0.0	0.3	27:01.
10	-10	128M	17252	14804	R	0.0	0.8	14:59.
10	-10	128M	17252	14804	S	0.0	0.8	2:33.
10	-10	128M	17252	14804	S	0.0	0.8	5:25.
20	0	77700	44800	3044	S	0.0	2.2	3:17.
20	0	44756	2868	1616	S	0.0	0.1	0:18.
20	0	28628	2068	1740	S	0.0	0.1	0:20.
20	0	29000	1280	1000	S	0.0	0.1	0:52.

# Linux Memory Management

Demo program 1: allocate memory without using it.

```
#include <stdio.h>
#include <stdlib.h>

int main (void) {
    int n = 0;

    while (1) {
        if (malloc(1<<20) == NULL) {
            printf("malloc failure after %d MiB\n", n)
            return 0;
        }
        printf ("got %d MiB\n", ++n);
    }
}
```

Demo program 2: allocate memory and actually touch it all.

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>

int main (void) {
    int n = 0;
    char *p;

    while (1) {
        if ((p = malloc(1<<20)) == NULL) {
            printf("malloc failure after %d MiB\n", n)
            return 0;
        }
        memset (p, 0, (1<<20));
        printf ("got %d MiB\n", ++n);
    }
}
```

<https://www.win.tue.nl/~aeb/linux/lk/lk-9.html>

Typically, the first demo program will get a very large amount of memory before `malloc()` returns `NULL`. The second demo program will get a much smaller amount of memory, now that earlier obtained memory is actually used.

# Linux Memory Management

Demo program 3: first allocate, and use later.

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>

#define N      10000

int main (void) {
    int i, n = 0;
    char *pp[N];

    for (n = 0; n < N; n++) {
        pp[n] = malloc(1<<20);
        if (pp[n] == NULL)
            break;
    }
    printf("malloc failure after %d MiB\n", n);

    for (i = 0; i < n; i++) {
        memset (pp[i], 0, (1<<20));
        printf("%d\n", i+1);
    }

    return 0;
}
```

- Allocate physical memory only if it is touched.

For example:

- On an 8 MiB machine without swap running 1.2.11:  
demo1: 274 MiB, demo2: 4 MiB, demo3: 270 / oom after 1 MiB: killed.
- Idem, with 32 MiB swap:  
demo1: 1528 MiB, demo2: 36 MiB, demo3: 1528 / oom after 23 MiB: killed.
- On a 32 MiB machine without swap running 2.0.34:  
demo1: 1919 MiB, demo2: 11 MiB, demo3: 1919 / oom after 4 MiB: bus error.
- Idem, with 62 MiB swap:  
demo1: 1919 MiB, demo2: 81 MiB, demo3: 1919 / oom after 74 MiB: The machine hangs. After several seconds: out of memory for bash. Out of memory for crond. Bus error.
- On a 256 MiB machine without swap running 2.6.8.1:  
demo1: 2933 MiB, demo2: after 98 MiB: killed. Also: Out of Memory: Killed process 17384 (java\_vm). demo3: 2933 / oom after 135 MiB: killed.
- Idem, with 539 MiB swap:  
demo1: 2933 MiB, demo2: after 635 MiB: killed. demo3: oom after 624 MiB: killed.

<https://www.win.tue.nl/~aeb/linux/lk/lk-9.html>

The third program will get the same large amount as the first program, and then is killed when it wants to use its memory. (On a well-functioning system, like Solaris, the three demo programs obtain the same amount of memory and do not crash but see `malloc()` return `NULL`.)

# Linux Memory Management

```
/proc/[pid]/status
Provides much of the information in /proc/[pid]/stat and /proc/[pid]/statm in a format that's easier for humans to parse. Here's an example:
```

```
$ cat /proc/$$/status
Name: bash
State: S (sleeping)
Tgid: 3515
Pid: 3515
PPid: 3452
TracerPid: 0
Uid: 1000 1000 1000 1000
Gid: 100 100 100 100
FDSize: 256
CpuUser: 16 00 100
CpuSystem: 00 00 000
VmPeak: 9136 kB
VmSize: 7896 kB
VmLck: 0 kB
VmPin: 0 kB
VmHWM: 7572 kB
VmRSS: 6316 kB
VmData: 5224 kB
VmStk: 88 kB
VmExe: 572 kB
VmLib: 1708 kB
VmPMD: 4 kB
VmPTE: 20 kB
VmSwap: 0 kB
Threads: 1
```

```
SigQ: 0/3067
SigPnd: 0000000000000000
ShdPnd: 0000000000000000
SigBlk: 000000000010000
SigIgn: 0000000000384004
SigCgt: 000000004b813efb
CapInh: 0000000000000000
CapPrm: 0000000000000000
CapEff: 0000000000000000
CapBnd: ffffffff
CapAmb: 0000000000000000
Seccomp: 0
Cpus_allowed: 00000001
Cpus_allowed_list: 0
Mems_allowed: 1
Mems_allowed_list: 0
voluntary_ctxt_switches: 150
nonvoluntary_ctxt_switches: 545
```

# Linux Memory Management

```
/proc/[pid]/maps
A file containing the currently mapped memory regions and their access permissions. See mmap(2) for some further information about memory mappings.

The format of the file is:

address      perms offset  dev  inode      pathname
00400000-00452000 r-xp 00000000 08:02 173521  /usr/bin/dbus-daemon
00651000-00652000 r--p 00051000 08:02 173521  /usr/bin/dbus-daemon
00652000-00655000 rw-p 00052000 08:02 173521  /usr/bin/dbus-daemon
00e03000-00e24000 rw-p 00000000 00:00 0       [heap]
00e24000-011f7000 rw-p 00000000 00:00 0       [heap]
...
35b1800000-35b1820000 r-xp 00000000 08:02 135522  /usr/lib64/ld-2.15.so
35b1a1f000-35b1a20000 r--p 0001f000 08:02 135522  /usr/lib64/ld-2.15.so
35b1a20000-35b1a21000 rw-p 00020000 08:02 135522  /usr/lib64/ld-2.15.so
35b1a21000-35b1a22000 rw-p 00000000 00:00 0
35b1c00000-35b1dac000 r-xp 00000000 08:02 135870  /usr/lib64/libc-2.15.so
35b1dac000-35b1fac000 ---p 001ac000 08:02 135870  /usr/lib64/libc-2.15.so
35b1fac000-35b1fb0000 r--p 001ac000 08:02 135870  /usr/lib64/libc-2.15.so
35b1fb0000-35b1fb2000 rw-p 001b0000 08:02 135870  /usr/lib64/libc-2.15.so
...
f2c6ff8c000-7f2c7078c000 rw-p 00000000 00:00 0       [stack:986]
...
7ffffb2c0d000-7ffffb2c2e000 rw-p 00000000 00:00 0       [stack]
7ffffb2d48000-7ffffb2d49000 r-xp 00000000 00:00 0       [vds]
```

The address field is the address space in the process that the mapping occupies. The perms field is a set of permissions:

```
r = read
w = write
x = execute
s = shared
p = private (copy on write)
```

The offset field is the offset into the file/whatever; dev is the device (major:minor); inode is the inode on that device. 0 indicates that no inode is associated with the memory region, as would be the case with BSS (uninitialized data).

The pathname field will usually be the file that is backing the mapping. For ELF files, you can easily coordinate with the offset field by looking at the Offset field in the ELF program headers (`readelf -l`).

There are additional helpful pseudo-paths:

[stack]  
The initial process's (also known as the main thread's) stack.

[stack:<tid>] (since Linux 3.4)  
A thread's stack (where the <tid> is a thread ID). It corresponds to the `/proc/[pid]/task/[tid]/` path.

[vds] The virtual dynamically linked shared object.

[heap] The process's heap.

If the pathname field is blank, this is an anonymous mapping as obtained via the `mmap(2)` function. There is no easy way to coordinate this back to a process's source, short of running it through `gdb(1)`, `strace(1)`, or similar.

Under Linux 2.0, there is no field giving pathname.

# Linux Memory Management

- `/proc/[pid]/mem`: This file can be used to access the pages of a process's memory through `open(2)`, `read(2)`, and `lseek(2)`.

```
# open and read mem
try:
    mem_file = open(mem_filename, 'rb+')
except IOError as e:
    print("[ERROR] Can not open file {}".format(mem_filename))
    print("    I/O error({}): {}".format(e.errno, e.strerror))
    maps_file.close()
    exit(1)

# read heap
mem_file.seek(addr_start)
heap = mem_file.read(addr_end - addr_start)

# find string
try:
    i = heap.index(bytes(search_string, "ASCII"))
except Exception:
    print("Can't find '{}'".format(search_string))
    maps_file.close()
    mem_file.close()
    exit(0)
print("[*] Found '{}' at {:x}".format(search_string, i))

# write the new string
print("[*] Writing '{}' at {:x}".format(write_string, addr_start + i))
mem_file.seek(addr_start + i)
mem_file.write(bytes(write_string, "ASCII"))

# close files
maps_file.close()
mem_file.close()

# there is only one heap in our example
break
```

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>

/**
 * main - uses strdup to create a new string, loops forever-ever
 *
 * Return: EXIT_FAILURE if malloc failed. Other never returns
 */
int main(void)
{
    char *s;
    unsigned long int i;

    s = strdup("Holberton");
    if (s == NULL)
    {
        fprintf(stderr, "Can't allocate mem with malloc\n");
        return (EXIT_FAILURE);
    }
    i = 0;
    while (s)
    {
        printf("[%lu] %s (%p)\n", i, s, (void *)s);
        sleep(1);
        i++;
    }
    return (EXIT_SUCCESS);
```

```
676] Holberton (0x10ff010)
677] Holberton (0x10ff010)
678] Holberton (0x10ff010)
679] Holberton (0x10ff010)
680] Holberton (0x10ff010)
681] Holberton (0x10ff010)
682] Fun w vm! (0x10ff010)
683] Fun w vm! (0x10ff010)
684] Fun w vm! (0x10ff010)
685] Fun w vm! (0x10ff010)
```

# **Memory Footprint of a Java process**

# Memory Footprint of a Java process

- Native Memory Tracking

```
-XX:NativeMemoryTracking=detail /  
summary  
-XX:+UnlockDiagnosticVMOptions  
-XX:+PrintNMTStatistics
```

NMT 中 committed，大致可以认为是 RSS 对应的值。

NMT 只能 Track 到 JVM 自身的内存分配情况，比如：Heap 内存分配，direct byte buffer 等；不能 track 到 jni 里直接调用 malloc 时的内存分配，这里最典型的就是 ZipStream 的场景。

一般情况下NMT Track出来的 committed 内存值既可能比RSS值大，也可能比RSS小，主要原因是：  
- 比真实RSS小： NMT 只能Track JVM自身的内

```
Total: reserved=3535846KB, committed=2261270KB
-
-           Java Heap (reserved=1257472KB, committed=1257472KB)
-                   (mmap: reserved=1257472KB, committed=1257472KB)

-
-               Class (reserved=1087227KB, committed=42491KB)
-                   (classes #6025)
-                   (malloc=5883KB #7385)
-                   (mmap: reserved=1081344KB, committed=36608KB)

-
-               Thread (reserved=39133KB, committed=39133KB)
-                   (thread #38)
-                   (stack: reserved=38964KB, committed=38964KB)
-                   (malloc=125KB #196)
-                   (arena=45KB #76)

-
-               Code (reserved=252732KB, committed=22892KB)
-                   (malloc=3132KB #5230)
-                   (mmap: reserved=249600KB, committed=19760KB)

-
-               GC (reserved=52265KB, committed=52265KB)
-                   (malloc=5773KB #151)
-                   (mmap: reserved=46492KB, committed=46492KB)

-
-               Compiler (reserved=172KB, committed=172KB)
-                   (malloc=41KB #225)
-                   (arena=131KB #3)

-
-               Internal (reserved=835091KB, committed=835091KB)
-                   (malloc=835059KB #15987)
-                   (mmap: reserved=32KB, committed=32KB)

-
-               Symbol (reserved=9562KB, committed=9562KB)
-                   (malloc=7891KB #74545)
-                   (arena=1670KB #1)

-
-       Native Memory Tracking (reserved=1813KB, committed=1813KB)
-                   (malloc=153KB #2414)
-                   (tracking overhead=1661KB)

-
-       Arena Chunk (reserved=378KB, committed=378KB)
-                   (malloc=378KB)
```

# Memory Footprint of a Java process

- Heap
  - Xms: Initial Heap size, NOT minimal
  - Xmx: Max Heap size
  - XX:+AlwaysPreTouch

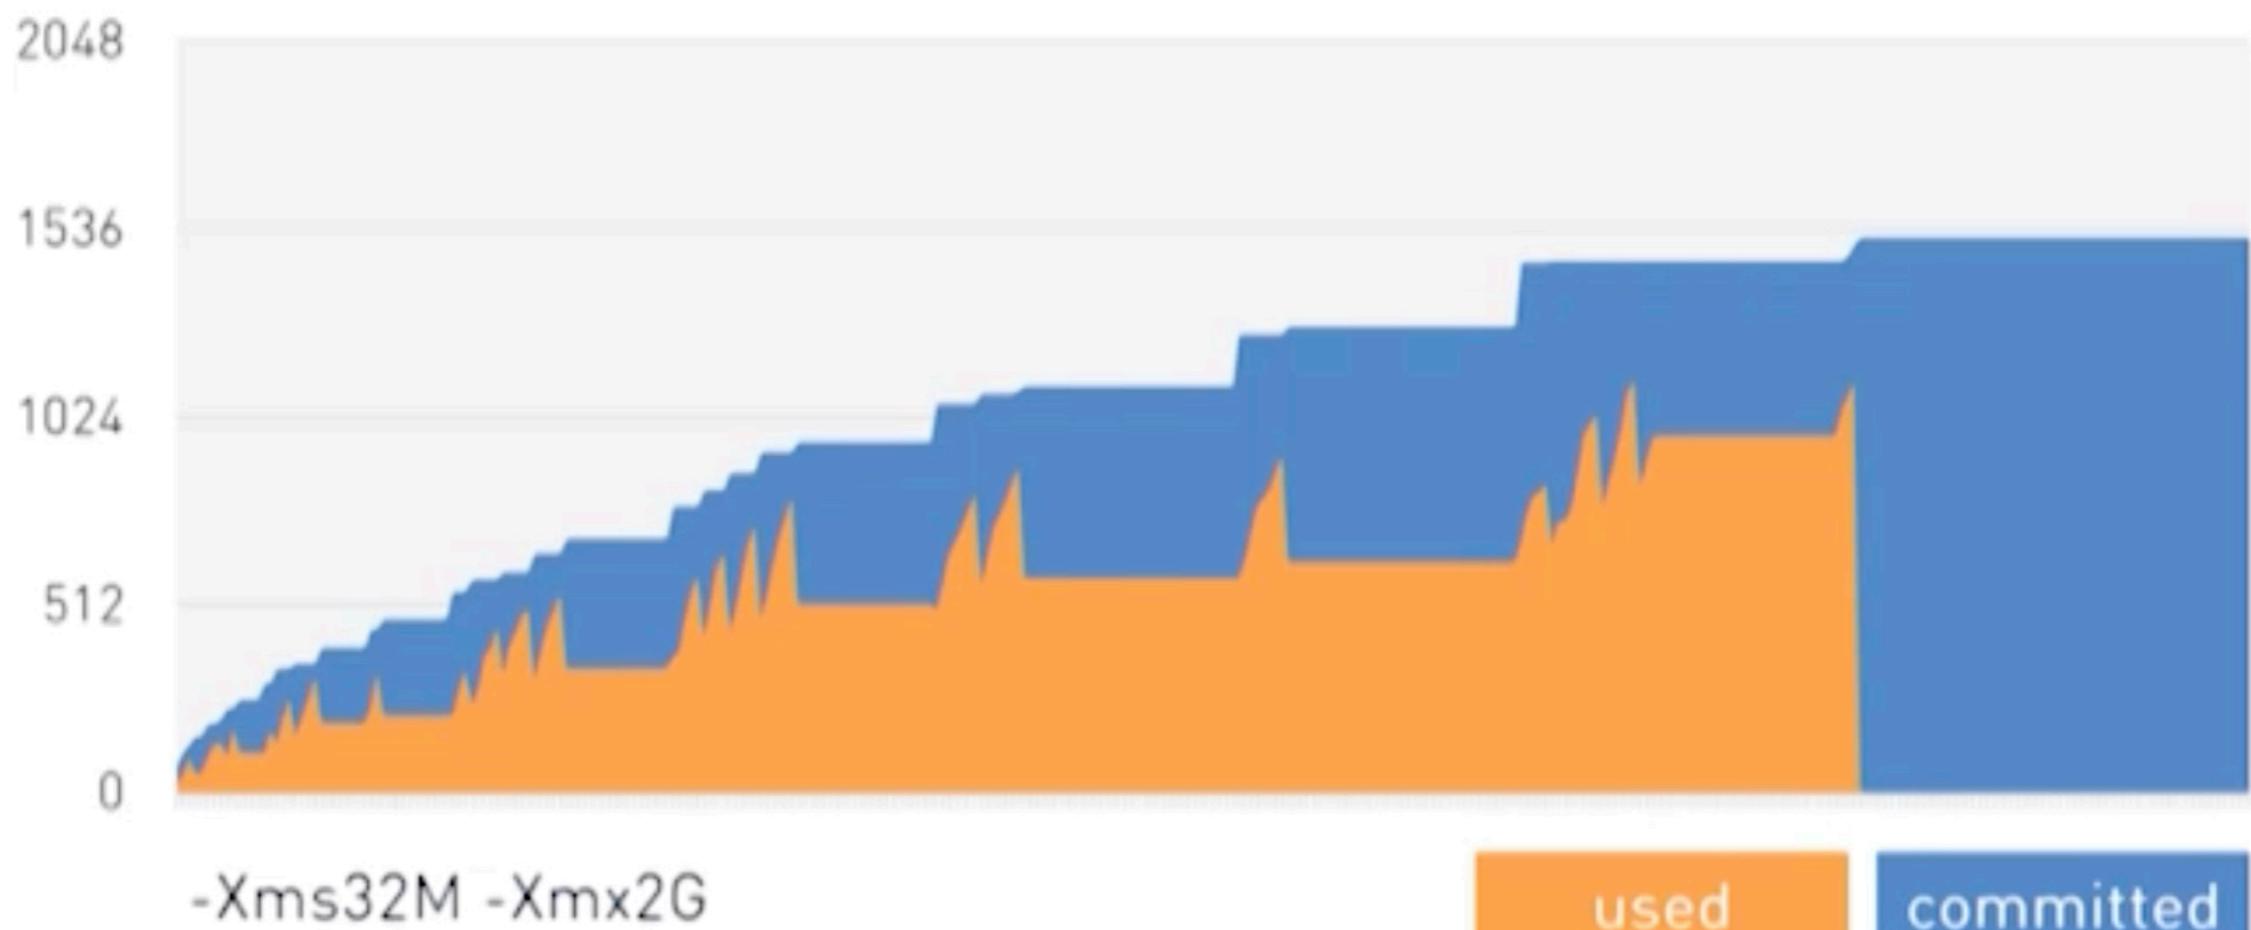
Pre-touch the Java heap during JVM initialization. Every page of the heap is thus demand-zeroed during initialization rather than incrementally during application execution. 【会】

# Memory Footprint of a Java process

- Heap
  - -XX:+UseAdaptiveSizePolicy(On by default)
  - -XX:MinHeapFreeRatio=40(expansion)
  - -XX:MaxHeapFreeRatio=70(shrinking)

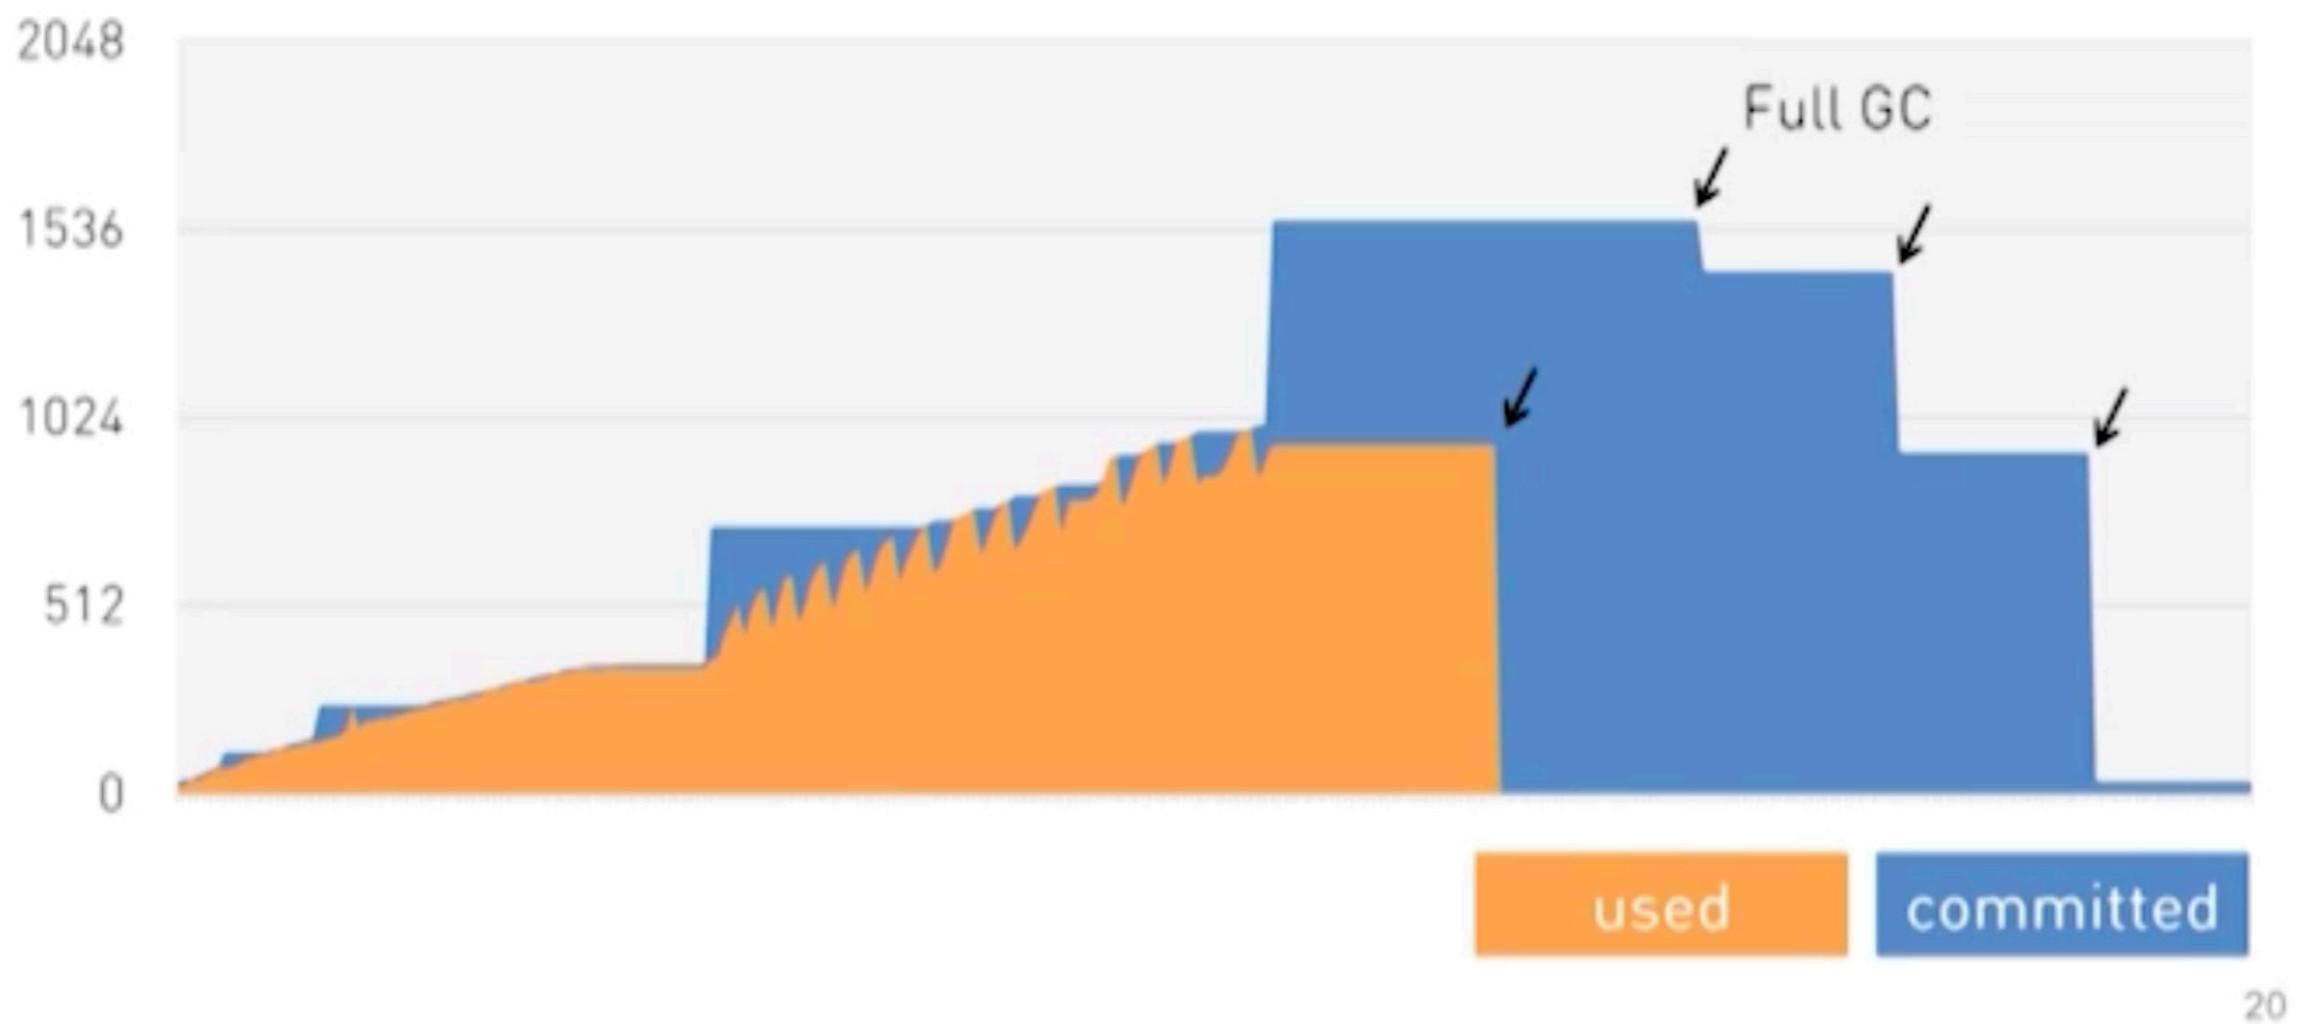


## Uncommit: Parallel





## Uncommit: CMS



## Uncommit: G1



# Memory Footprint of a Java process

- Metaspace:
  - Class: More classes are loaded, more metaspace is used.
  - Methods
  - Annotations
  - ....
- JVM Opt:
  - -XX:MaxMetaspaceSize (unlimited by default)
  - -XX:MetaspaceSize=64M

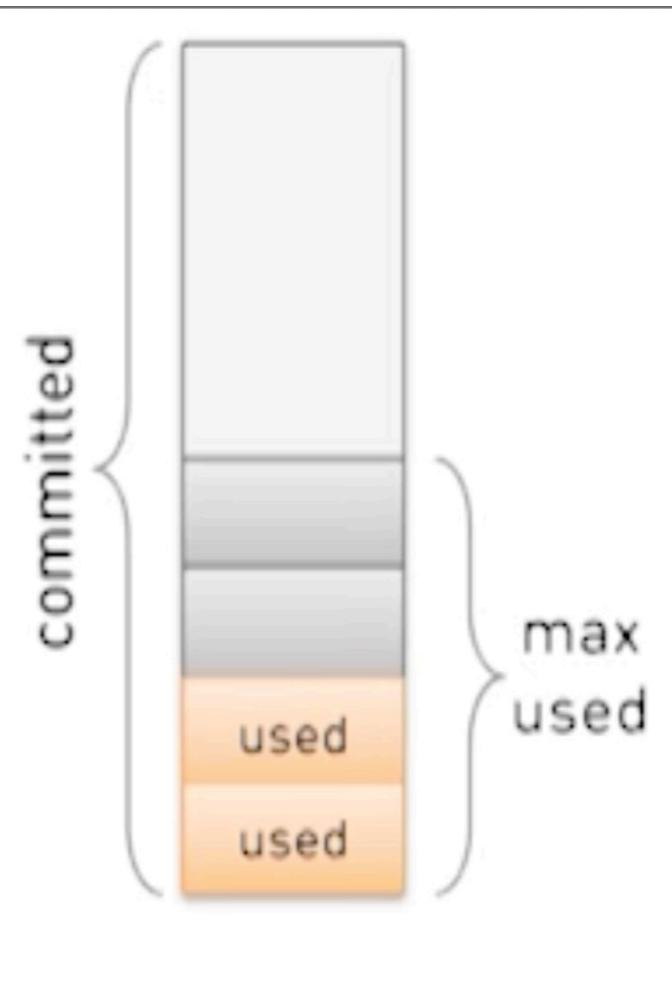
Class metadata (method bytecodes, symbols, constant pools, annotations etc.) is stored in off-heap area called Metaspace. The more classes are loaded - the more metaspace is used.

```
2400.241: [GC concurrent-root-region-scan-start]
2400.241: [Full GC (Metadata GC Threshold) 2400.252: [GC concurrent-root-region-sc
2400.252: [GC concurrent-mark-start]
1151M->603M(4356M), 2.6980537 secs]
[Eden: 0.0B(2558.0M)->0.0B(2613.0M) Survivors: 55.0M->0.0B Heap: 1151.7M(4356.0
[Times: user=3.92 sys=0.00, real=2.70 secs]
```

# Memory Footprint of a Java process

```
- Thread (reserved=39133KB, committed=39133KB)
      (thread #38)
      (stack: reserved=38964KB, committed=38964KB)
      (malloc=125KB #196)
      (arena=45KB #76)
```

- Thread
  - Xss
  - Committed!=Resident



# Memory Footprint of a Java process

## Native Memory Tracking:

```
Total: reserved=4023326KB committed=2762382KB
- Java Heap (reserved=1331200KB committed=1331200KB)
  (mmap: reserved=1331200KB, committed=1331200KB)

- Class (reserved=1108143KB, committed=64559KB)
  (classes #8621)
  (malloc=6319KB #17371)
  (mmap: reserved=1101824KB, committed=58240KB)

- Thread (reserved=1100668KB committed=1190668KB)
  (thread #1154)
  (stack: reserved=1185284KB, committed=1185284KB)
  (malloc=3809KB #5771)
  (arena=1575KB #2306)

- Code (reserved=255744KB, committed=38384KB)
  (malloc=6144KB #8858)
  (mmap: reserved=249600KB, committed=32240KB)

- GC (reserved=54995KB, committed=54995KB)
  (malloc=5775KB #217)
  (mmap: reserved=49220KB, committed=49220KB)

- Compiler (reserved=267KB, committed=267KB)
  (malloc=137KB #333)
  (arena=131KB #3)

- Internal (reserved=65106KB, committed=65106KB)
  (malloc=65074KB #29652)
  (mmap: reserved=32KB, committed=32KB)

- Symbol (reserved=13622KB, committed=13622KB)
  (malloc=12016KB #128199)
  (arena=1606KB #1)
```

This shows 2.7GB of committed memory, including 1.3GB of allocated heap and almost 1.2GB of allocated thread stacks (using many threads).

However, when running `ps ax -o pid,rss | grep <mypid>` or `top` it shows only 1.6GB of `RES/rss` resident memory. Checking swap says none in use:

	total	used	free	shared	buffers	cached
Mem:	129180	99348	29831	0	2689	73024
-/+ buffers/cache:		23633	105546			
Swap:	15624	0	15624			

Why does the JVM indicate 2.7GB memory is committed when only 1.6GB is resident? Where did the rest go?

# Memory Footprint of a Java process

- Code:
  - Contains dynamically generated code: JIT-compiled methods, interpreter and run-time stubs. Its size is limited by -  
`XX:ReservedCodeCacheSize (240M by default).`
  - Turn off `-XX:-TieredCompilation` to reduce the amount of compiled code and thus the Code Cache usage.
  - `-XX:+PrintCodeCache`

```
CodeCache: size=245760Kb used=1108Kb max_used=1108Kb free=244651Kb
bounds [0x0000000114d8e000, 0x0000000114ffe000, 0x0000000123d8e000]
total_blobs=269 nmethods=28 adapters=155
compilation: enabled
```

```
Process finished with exit code 0
```

# Memory Footprint of a Java process

- Garbage Collector:

- Mark Bitmap
- Mark Stack (for traversing object graph)

- Remembered Set

GC memory overhead

GC	Overhead, MB	%
Serial	7	0,3%
Shenandoah	38	1,9%
CMS	76	3,7%
Parallel	90	4,4%
G1	166	8,1%
Z	238	11,6%

OpenJDK 13 | Heap size = 2 GB

# Memory Footprint of a Java process

- Compiler:
  - JIT compiler itself also requires memory to do its job. This can be reduced again by switching off Tiered Compilation or by reducing the number of compiler threads: -XX:CICompilerCount.

# Memory Footprint of a Java process

- Internal

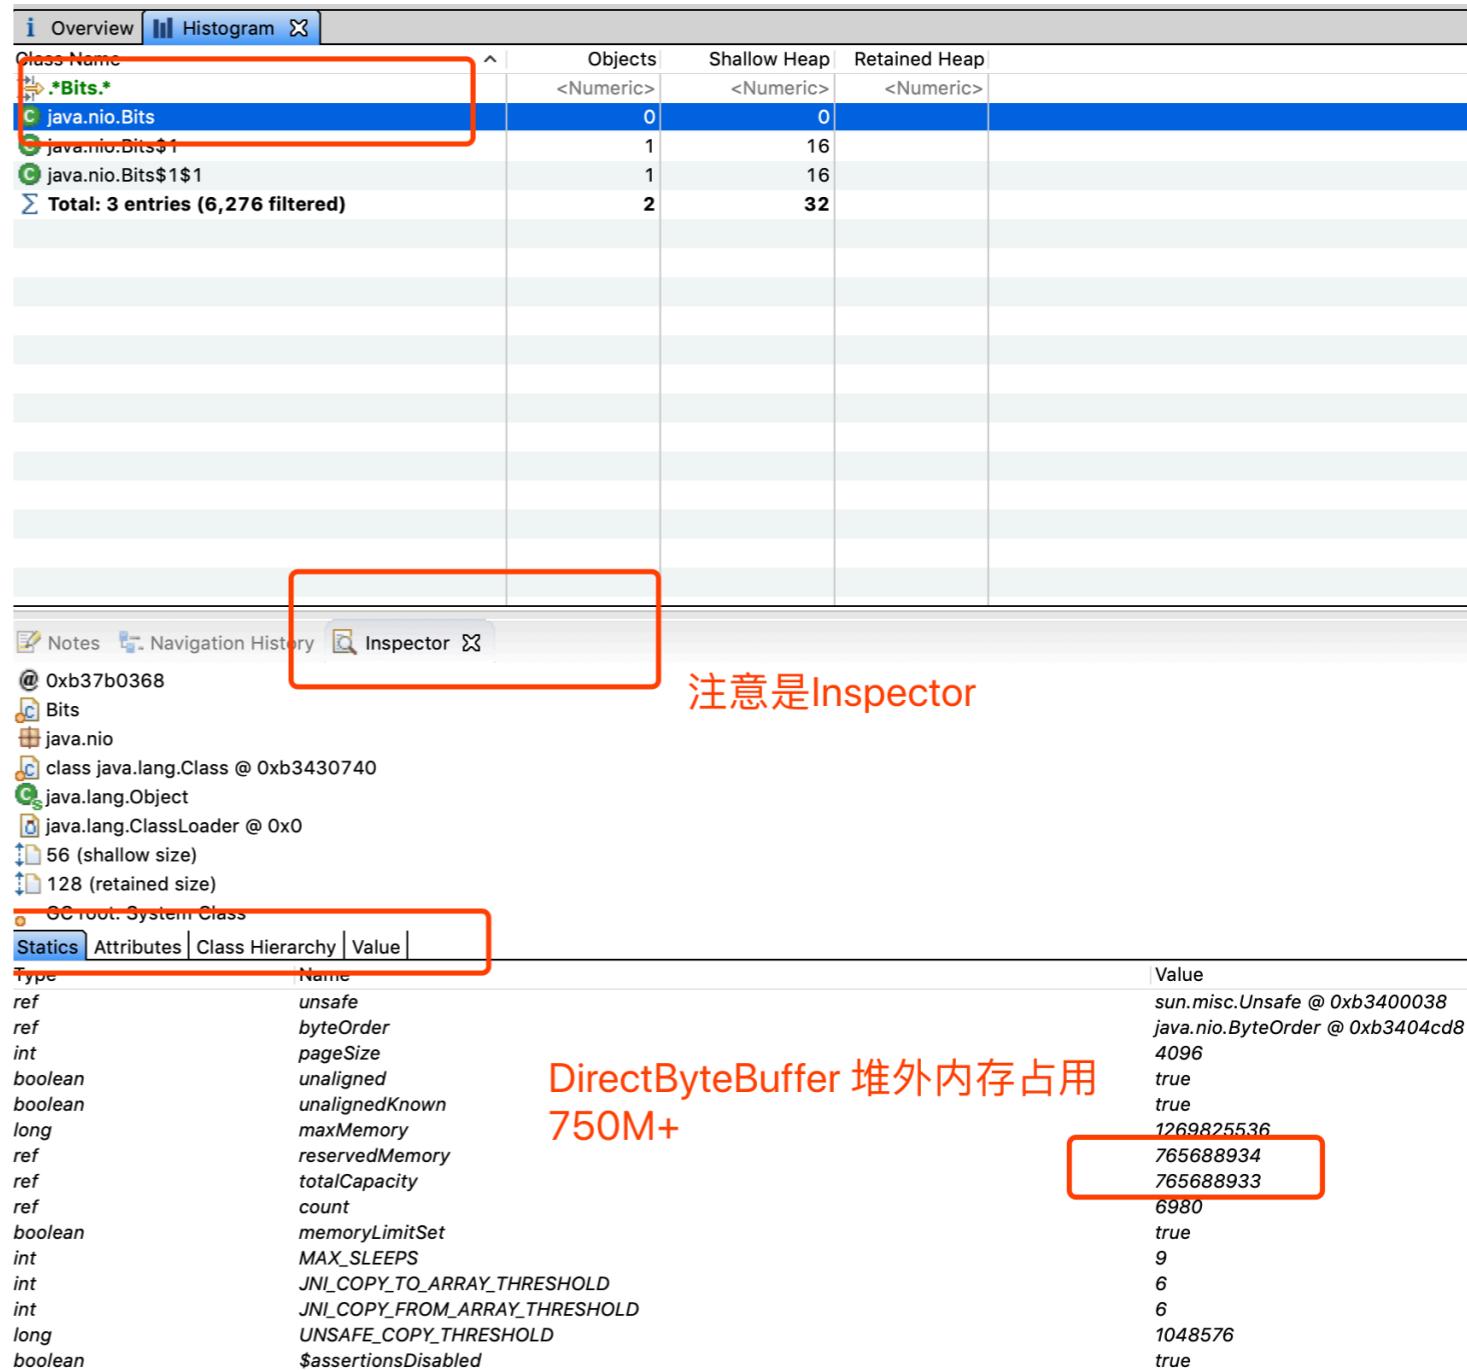
```
200 |                               (arena=151KB #3)
201 |
202 -           Internal (reserved=835091KB, committed=835091KB)
203 |               (malloc=835059KB #15987)
204 |               (mmap: reserved=32KB, committed=32KB)
205 |
206 -           Symbol (reserved=9562KB, committed=9562KB)
207 |               (malloc=7891KB #74545)
208 |               (arena=1670KB #1)
209 |
```

- Direct Byte Buffer

```
1374 |     [0x00007f6cdd49697e] JavaMain+0x9e
1375 |     Details:
1376 |
1377 |     [0x00007f6cdc95f03] Unsafe_AllocateMemory+0xc3
1378 |     [0x00007f6cc5f5c38]
1379 |         (malloc=800998KB #7495)
1380 |
1381 |     [0x00007f6cdc95f03] Unsafe_AllocateMemory+0xc3
1382 |     [0x00007f6cc501867]
1383 |         (malloc=27882KB #2447)
1384 |
1385 |     [0x00007f6cdc8574e5] ArrayAllocator<StarTask, (MemoryType)1>::allocate(unsigned long
1386 |     [0x00007f6cdc855c10] PSPromotionManager::PSPromotionManager()+0x1f0
1387 |     [0x00007f6cdc855e5d] PSPromotionManager::initialize()+0x13d
1388 |     [0x00007f6cdc95a724] universe_post_init()+0x8a4
1389 |         (malloc=5120KB #5)
1390 |
```

# Memory Footprint of a Java process

- Internal
  - Direct Byte Buffer



# Memory Footprint of a Java process

- Symbol:

```
- Symbol (reserved=1464KB, committed=1464KB)
          (malloc=944KB #161)
          (arena=520KB #1)

SymbolTable statistics:
Number of buckets      : 20011 = 160088 bytes, avg 8.000
Number of entries       : 12431 = 298344 bytes, avg 24.000
Number of literals      : 12431 = 477440 bytes, avg 38.407
Total footprint         :           = 935872 bytes
Average bucket size    : 0.621
Variance of bucket size: 0.620
Std. dev. of bucket size: 0.787
Maximum bucket size    :           6

StringTable statistics:
Number of buckets      : 60013 = 480104 bytes, avg 8.000
Number of entries       : 877 = 21048 bytes, avg 24.000
Number of literals      : 877 = 59120 bytes, avg 67.412
Total footprint         :           = 560272 bytes
Average bucket size    : 0.015
Variance of bucket size: 0.015
Std. dev. of bucket size: 0.121
Maximum bucket size    :           2
```

- SymbolTable: names signatures
- StringTable: interned strings
- -XX:+PrintStringTableStatistics

# Memory Footprint of a Java process

- Others
  - JNI malloc
  - ZipInputStream
  - .....

# Memory Footprint of a Java process

- Why is Java using much more memory than heap size?
  - Off-Heap
- How to keep the memory used by a Java process under control?
  - No way!
- How to detect the memory usage of off-heap?
  - pmap, strace,gdb....
- Why does NMT (Native Memory Tracking) report more / less committed memory than the linux process resident set size?
  - Less: JNI malloc
  - More: Thread (Committed != Resident)
- .....

# Tools to detect memory usage of a Java process

# Tools to detect memory usage of a Java process

- Heap
  - -XX:+PrintFlagsFinal
  - jinfo
  - jmap
  - MAT
  - .....

# Tools to detect memory usage of a Java process

- Off-Heap
  - NMT
  - pmap
  - gdb
  - strace
  - async-profiler
  - .....

# Examples

# Examples

- <https://zhuanlan.zhihu.com/p/54048271>
- <https://zhuanlan.zhihu.com/p/60976273>

**Thanks!**